

Feature Representation by Convolution


A convolutional layer is not fully connected, but has a narrowed down receptive field (e.g. 3×3).

The parameters of each filter are spatially shared: A feature that is useful in one place, ought to be useful in another, too.

Downsampling

Upsampling

Parameter sharing



- **Depth:** number of filters
- **Stride:** filter step size (when we “slide” it)
- **Padding:** zero-pad the input

<https://selfdrivingcars.mit.edu/>

Neural Network Segmentation - Basic structure

Feature Representation by Convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Zero-Padding is adding zero-valued pixel to the image border (gray area).

Downsampling

Upsampling

Parameter sharing

Neural Network Segmentation - Basic structure

Feature Representation by Convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Zero-padded image

| | | |
|----|----|----|
| 0 | 0 | -1 |
| -1 | 0 | 0 |
| -1 | -1 | -1 |

Filter

| |
|---|
| 0 |
|---|

Bias

| | | |
|----|----|----|
| -4 | -4 | 0 |
| -3 | -4 | -3 |
| 0 | -3 | -1 |

Output

Downsampling

Upsampling

Parameter sharing

Neural Network Segmentation - Basic structure

Feature Representation by Convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | |
|----|----|----|
| 0 | 0 | -1 |
| -1 | 0 | 0 |
| -1 | -1 | -1 |

Filter

| |
|---|
| 0 |
|---|

Bias

| | | |
|----|----|----|
| -4 | -4 | 0 |
| -3 | -4 | -3 |
| 0 | -3 | -1 |

Output

Amount of filters or convolution depth: 1

Filter step size or Stride: 2

Zero-padded image

Downsampling

Upsampling

Parameter sharing

Neural Network Segmentation - Basic structure

Feature Representation by Convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | |
|----|----|----|
| 0 | 0 | -1 |
| -1 | 0 | 0 |
| -1 | -1 | -1 |

Filter

| |
|---|
| 0 |
|---|

Bias

| | | |
|----|----|----|
| -4 | -4 | 0 |
| -3 | -4 | -3 |
| 0 | -3 | -1 |

Output

Review edge detector:

Similar idea, now the parameters of the filters are learned. We want a lot of filters!

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

Downsampling

Upsampling

Parameter sharing

Lena Test image

Lena , the 'hello world!' of image processing. 330x330

Cover photo of 1972 Playboy magazine of the Swedish model Lena Söderberg.

Since then Lena was a guest at several IEEE conferences. The image also sparked discussions on gender-equality in the male-dominated field of engineering.

It is a good test image because of its detail, flat regions, shading, and texture.



<https://en.wikipedia.org/wiki/Lenna>

Introduction – Conventional segmentation approaches

Thresholding

Edge detection

Segment boundaries and edges are closely related.

Since there is often a large gradient at the segment boundaries.



Canny Edge Detection

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

Sobel Operators for Edge Detection

Clustering

Region growing

Neural Network Segmentation - Basic structure

Feature Representation by Convolution

| | | | | | | |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Downsampling
Upsampling
Parameter sharing

| | | |
|----|----|----|
| 0 | 0 | -1 |
| -1 | 0 | 0 |
| -1 | -1 | -1 |

Filter

| |
|---|
| 0 |
|---|


Bias

| | | |
|----|----|----|
| -4 | -4 | 0 |
| -3 | -4 | -3 |
| 0 | -3 | -1 |

Output

Review edge detector:

Similar idea, now the parameters of the filters are learned. And we want to go deep!

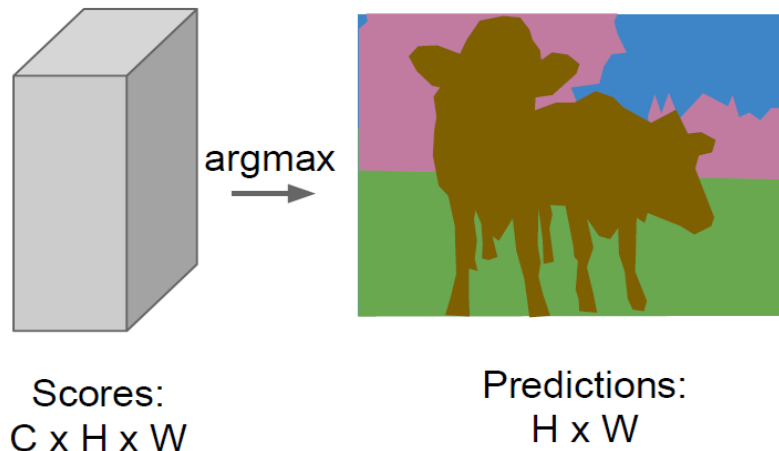


Neural Network Segmentation - Optimization

Last layer

Last Layer results in a tensor with $H \times W$ image resolution and a depth of C : Number of classes to segment.

The last layer should encode the values into a range of values of (0;1).
Either by **softmax** or **sigmoid** function.



<http://cs231n.github.io/>

Cross-entropy
Dice-coefficient

Neural Network Activation Functions - Review


Sigmoid

$$f(z) = \frac{1}{1 + e^{-z}}$$

$$f'(z) = f(z)(1 - f(z))$$

Binary classification only.

The probability sum does not need to be one.



<http://cs231n.github.io/neural-networks-1/>

Neural Network Activation Functions - Review

Softmax

Normalized exponential function

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Used for multi-class segmentation.

Probability sum will be 1.

before after

| | |
|-----|-----|
| 2.0 | 0.7 |
| 1.0 | 0.2 |
| 0.1 | 0.1 |